

Assignment 1 – Big Data Processing – Olivier Salaün – 16/02/2017.

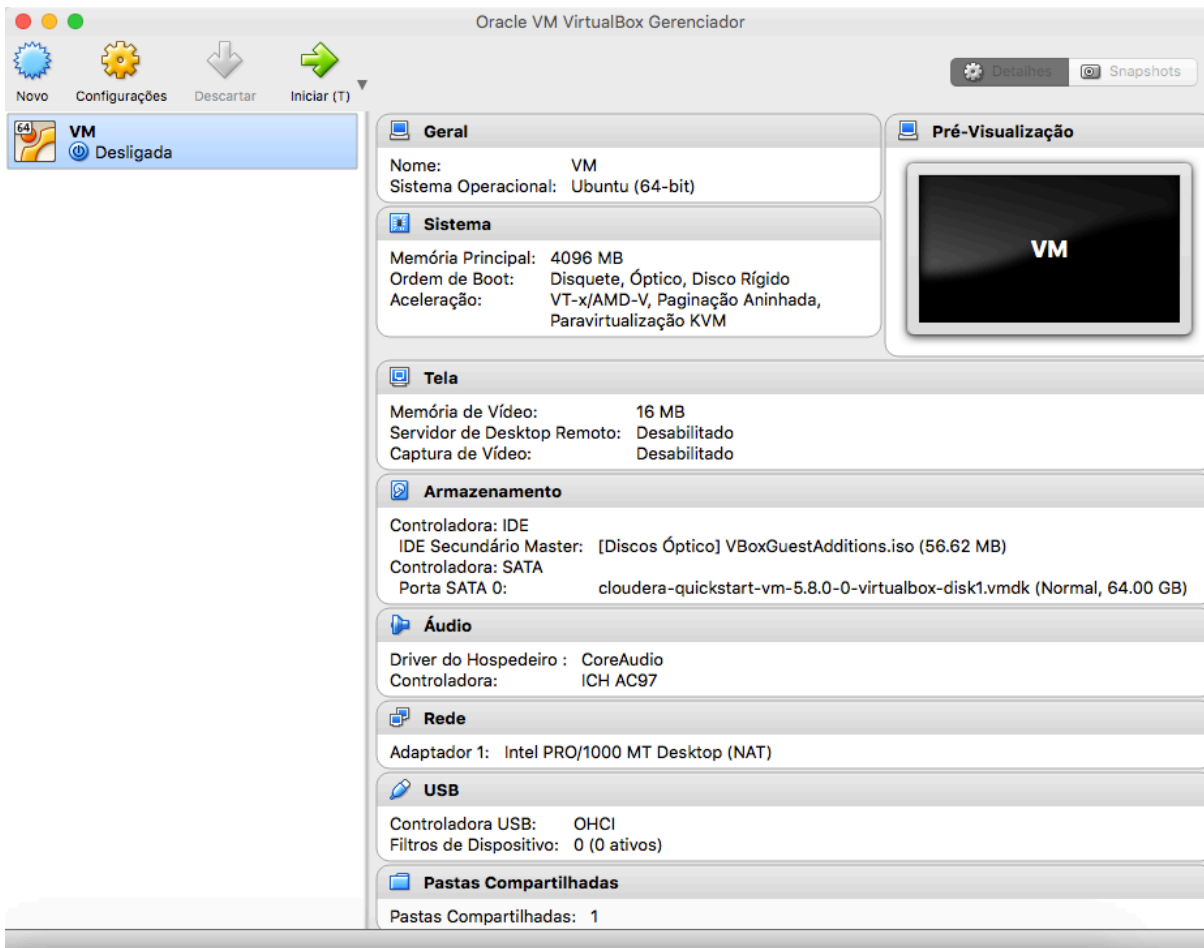
Settings.

The virtual machine is the same as the one used for the WordCount assignment. It is a CentOS that relies on a 64-bit version of Ubuntu with 4GB of allocated RAM and a 2.5 GHz CPU.

```
[cloudera@quickstart ~]$ hadoop version
Hadoop 2.6.0-cdh5.8.0
Subversion http://github.com/cloudera/hadoop -r
57e7b8556919574d517e874abfb7ebe31a366c2b
Compiled by jenkins on 2016-06-16T19:38Z
Compiled with protoc 2.5.0
From source with checksum 9e99ecd28376acfd5f78c325dd939fed
This command was run using /usr/lib/hadoop/hadoop-common-2.6.0-cdh5.8.0.jar
```

```
[cloudera@quickstart ~]$ lsb_release -a
LSB Version:          :base-4.0-amd64:base-4.0-noarch:core-4.0-amd64:core-4.0-noarch
Distributor ID:       CentOS
Description:          CentOS release 6.7 (Final)
Release:              6.7
Codename:             Final
```

```
[cloudera@quickstart ~]$ cat /proc/cpuinfo
processor       : 0
vendor_id     : GenuineIntel
cpu family    : 6
model        : 58
model name    : Intel(R) Core(TM) i5-3210M CPU @ 2.50GHz
stepping     : 9
microcode    : 25
cpu MHz      : 2494.316
cache size   : 3072 KB
physical id  : 0
siblings     : 1
core id      : 0
cpu cores    : 1
apicid       : 0
initial apicid : 0
fpu          : yes
fpu_exception : yes
cpuid level  : 13
wp           : yes
flags        : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat
pse36 clflush mmx fxsr sse sse2 syscall nx rdtscp lm constant_tsc up rep_good
xtopology nonstop_tsc unfair_spinlock pni pclmulqdq monitor ssse3 cx16 sse4_1
sse4_2 x2apic popcnt aes xsave avx rdrand hypervisor lahf_lm
bogomips     : 4988.63
clflush size  : 64
cache_alignment : 64
address sizes : 36 bits physical, 48 bits virtual
power management:
```



In order to keep things clear, we create a whole new folder for each question in Eclipse Package Explorer.

(a) (30) Run a MapReduce program to identify stop words (words with frequency > 4000) for the given document corpus. Store them in a single csv file on HDFS (stopwords.csv). You can edit the several parts of the reducers' output after the job finishes (with hdfs commands or with a text editor), in order to merge them as a single csv file.

For all questions a), the (key, values) formats are the following:

mapper input: (LongWritable, Text) [file names and content of the files]

mapper output – reducer input: (Text, IntWritable) [word, number of occurrences of that word]

reducer output: (Text, IntWritable) [word, number of occurrences of that word]

i. (10) Use 10 reducers and do not use a combiner. Report the execution time.

We create an input folder where all three texts are initially stored and we reuse the code from the WordCount assignment for this question. In the driver class, we add 10 reducers with the following line :

```
job.setNumReduceTasks(10);
```

We also add an if-condition so that existing arguments files are replaced by new ones when the program is executed repeatedly.

```
FileSystem fs = FileSystem.newInstance(getConf());
```

```
if (fs.exists(new Path(args[1]))) {  
    fs.delete(new Path(args[1]), true);  
}
```

Since the output is a csv file, we use semicolon as a separator:

```
job.getConfiguration().set("mapreduce.output.textoutputformat.separator", " ; ");
```

In the map class, we use a for-loop that splits the text at whitespaces between words and convert all words into lowercase strings.

```
for (String string: value.toString().split("\\s+")) {  
    word.set(string.toLowerCase());  
    context.write(word, ONE);  
}
```

In the reduce class, we create another loop that generates a counter for each word and those that have more than 4000 occurrences are included within the output file.

```
for (IntWritable val : values) {  
    sum = sum + val.get();  
}  
if (sum > 4000){
```

In the outputfile, each key (stop word) comes with its value (integer, number of occurrences):

```
context.write(key, new IntWritable(sum));  
}
```

We use the following commands in the terminal :

```
cd workspace  
hdfs dfsadmin -safemode leave  
hadoop jar InvertedIndex_a_i.jar stopwordcount.InvertedIndex_a_i  
input output  
hadoop fs -ls output  
hadoop fs -getmerge output  
/home/cloudera/workspace/InvertedIndex_a_i/output/StopWords_10redu  
c_nocomb.csv
```

The whole process takes 3minutes and 47seconds.

The first screenshot shows the 'Application Overview' for 'InvertedIndex_a_i'. The application is in a 'FINISHED' state, succeeded, and took 3 minutes and 47 seconds to complete. The application metrics show zero preempted resources and an aggregate allocation of 1461124 MB-seconds and 1186 vcore-seconds. The application master table shows one attempt starting at 12:16:47 on node quickstart.cloudera:8042.

The second screenshot shows the 'Job Overview' for the same application. The job is in a 'SUCCEEDED' state, took 3 minutes and 16 seconds to complete, and was submitted at 12:16:46 HST. The job diagnostics show an average map time of 1 minute and 12 seconds, an average shuffle time of 48 seconds, an average merge time of 6 seconds, and an average reduce time of 11 seconds. The application master table shows one attempt starting at 12:17:00 HST on node quickstart.cloudera:8042. The task summary table shows 3 map tasks and 10 reduce tasks, all of which were successful.

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:16:47 -1000 2017	quickstart.cloudera:8042	logs

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:17:00 HST 2017	quickstart.cloudera:8042	logs

Task Type	Total	Complete
Map	3	3
Reduce	10	10

Attempt Type	Failed	Killed	Successful
Maps	0	0	3
Reduces	0	3	10

ii. (10) Run the same program again, this time using a Combiner. Report the execution time. Is there any difference in the execution time, compared to the previous execution? Why?

The purpose of a combiner is to take the mapper output and to create a reducer input that is minimized. We take the code used previously and add a combiner in the driver class with this line :

```
job.setCombinerClass(Reduce.class);
```

The whole process takes 2 minutes and 52 seconds, showing that the combiner allows a 24% reduction in execution runtime with respect to previous question thanks to minimized reducer input.

The first screenshot shows the 'Application Overview' for a MapReduce application. The application is named 'InvertedIndex_a_ii', is in a 'FINISHED' state, and has a 'FinalStatus' of 'SUCCEEDED'. It was submitted on 'Wed Feb 15 12:21:15 -1000 2017' and took '2mins, 52sec' to complete. The 'Tracking URL' is 'History'.

The second screenshot shows the 'Job Overview' for the same application. The job is named 'InvertedIndex_a_ii', is in a 'SUCCEEDED' state, and was submitted on 'Wed Feb 15 12:21:15 HST 2017'. It took '2mins, 29sec' to complete. The 'Average Map Time' is 59sec, 'Average Shuffle Time' is 42sec, 'Average Merge Time' is 0sec, and 'Average Reduce Time' is 5sec.

Both screenshots include a table for 'ApplicationMaster' attempts, showing the first attempt on 'Wed Feb 15 12:21:15 -1000 2017' on node 'quickstart.cloudera:8042'.

iii. (5) Run the same program again, this time compressing the intermediate results of map (using any codec you wish). Report the execution time. Is there any difference in the execution, time compared to the previous execution? Why?

The purpose of a compressor is to compress input data so that they are faster to read during the hadoop process. It also helps in reducing storage space usage.

We import the following compressors in the code used previously:

```
import org.apache.hadoop.io.compress.BZip2Codec;
import org.apache.hadoop.io.compress.CompressionCodec;
```

The execution time is 2 minutes and 54 seconds, we got almost the same performance like in question a)ii). This might be due to the fact that compression requires significant amount of time. Some codecs do not allow splitting of compressed files but bzip2 supports splitting, so it should not be an issue in our case.

The image contains two screenshots of the Cloudera interface, likely Cloudera Desktop or Cloudera Manager, showing details for a specific application and its corresponding MapReduce job.

Top Screenshot: Application Overview

The top screenshot shows the "Application Overview" page for the application "InvertedIndex_a_iii". The application is in a "FINISHED" state. The "FinalStatus" is "SUCCEEDED". The "Started" time is "Wed Feb 15 12:24:48 -1000 2017". The "Elapsed" time is "2mins, 54sec". The "Tracking URL" is "History". The "Diagnostics" section is empty.

Application Metrics

Total Resource Preempted:	<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:	0
Total Number of AM Containers Preempted:	0
Resource Preempted from Current Attempt:	<memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt:	0
Aggregate Resource Allocation:	1129295 MB-seconds, 914 vcore-seconds

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:24:48 -1000 2017	quickstart.cloudera:8042	logs

Bottom Screenshot: Job Overview

The bottom screenshot shows the "Job Overview" page for the job "InvertedIndex_a_iii". The job is in a "SUCCEEDED" state. The "Submitted" time is "Wed Feb 15 12:24:48 HST 2017". The "Started" time is "Wed Feb 15 12:25:11 HST 2017". The "Finished" time is "Wed Feb 15 12:27:43 HST 2017". The "Elapsed" time is "2mins, 32sec". The "Diagnostics" section is empty.

Job Metrics

Average Map Time	59sec
Average Shuffle Time	43sec
Average Merge Time	0sec
Average Reduce Time	6sec

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:24:54 HST 2017	quickstart.cloudera:8042	logs

Task Summary

Task Type	Total	Complete
Map	3	3
Reduce	10	10

Attempt Summary

Attempt Type	Failed	Killed	Successful
Maps	0	0	3
Reduces	0	2	10

iv. (5) Run the same program again, this time using 50 reducers. Report the execution time. Is there any difference in the execution time, compared to the previous execution? Why?

The purpose of a reducer is to reduce a set of intermediate values corresponding to a common key into a smaller set. Augmenting the number of reducers allows running many reducing tasks in parallel, but the number of reducers has to be a balance:

- if there are too few reducers with respect to the amount of data to process, it implies that more tasks could have been parallelized. The job will be slow with the few existing reducers;
- if the reducers are too many, many tasks will be executed in parallel on very small amount of files, that is time consuming and not optimal.

Unsurprisingly, in this question, setting the number of reducers at 50 causes a dramatic rise in the execution time: 8 minutes and 15 seconds. We used the following line in the driver class:

```
job.setNumReduceTasks(50);
```

We get a 184% increase with respect to the execution time obtained in previous question. There are clearly too many reducers with respect to the amount of files to process.

The image contains two screenshots from the Cloudera interface, likely Cloudera Manager, showing details for a MapReduce application and its job.

Top Screenshot: Application Overview

The browser address bar shows `http://localhost:8088/cluster/app/application_1487196663434_0004`. The left sidebar shows a navigation menu with options like Cluster, About, Nodes, Applications, NEW, NEW_SAVING, SUBMITTED, ACCEPTED, RUNNING, FINISHED, FAILED, KILLED, and Scheduler. The main content area displays the following details:

- User:** cloudera
- Name:** InvertedIndex_a_iv
- Application Type:** MAPREDUCE
- Application Tags:**
- State:** FINISHED
- FinalStatus:** SUCCEEDED
- Started:** Wed Feb 15 12:28:22 -1000 2017
- Elapsed:** 8mins, 15sec
- Tracking URL:** History
- Diagnostics:**

Below this, the **Application Metrics** section shows:

- Total Resource Preempted:** <memory:0, vCores:0>
- Total Number of Non-AM Containers Preempted:** 0
- Total Number of AM Containers Preempted:** 0
- Resource Preempted from Current Attempt:** <memory:0, vCores:0>
- Number of Non-AM Containers Preempted from Current Attempt:** 0
- Aggregate Resource Allocation:** 3690663 MB-seconds, 3078 vcore-seconds

The **ApplicationMaster** table shows one attempt:

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:28:22 -1000 2017	quickstart.cloudera:8042	logs

Bottom Screenshot: MapReduce Job Overview

The browser address bar shows `quickstart.cloudera:19888/jobhistory/job/job_1487196663434_0004/jobhistory/job/job_1487196`. The left sidebar shows a navigation menu with options like Application, Job, Overview, Counters, Configuration, Map tasks, Reduce tasks, and Tools. The main content area displays the following details:

- Job Name:** InvertedIndex_a_iv
- User Name:** cloudera
- Queue:** root.cloudera
- State:** SUCCEEDED
- Uberized:** false
- Submitted:** Wed Feb 15 12:28:22 HST 2017
- Started:** Wed Feb 15 12:28:45 HST 2017
- Finished:** Wed Feb 15 12:36:38 HST 2017
- Elapsed:** 7mins, 52sec
- Diagnostics:**
- Average Map Time:** 58sec
- Average Shuffle Time:** 42sec
- Average Merge Time:** 0sec
- Average Reduce Time:** 3sec

The **ApplicationMaster** table shows one attempt:

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:28:28 HST 2017	quickstart.cloudera:8042	logs

Below this, the **Task Type** table shows the distribution of tasks:

Task Type	Total	Complete
Map	3	3
Reduce	50	50

The **Attempt Type** table shows the status of tasks:

Attempt Type	Failed	Killed	Successful
Maps	0	0	3
Reduces	0	2	50

(b) (30) Implement a simple inverted index for the given document corpus, as shown in the previous Table, skipping the words of stopwords.csv.

The (key, values) formats are the following:

mapper input: (LongWritable, Text) [file names and content of the files]

mapper output – reducer input: (Text, Text) [word, file name where the word appears]

reducer output: (Text, Text) [word, file name where the word appears]

We reuse code from previous questions and we set the number of reducers at 10 in the driver class.

```
job.setNumReduceTasks(10);
```

Our folder contains the three text input files and a StopWords.csv obtained from questions a). In the mapper class, we create a HashSet for storing the stop words from the csv file.

```
HashSet<String> stopwordslist = new HashSet<String>();
    BufferedReader Reader = new BufferedReader(
        new FileReader(new
            File("/home/cloudera/workspace/InvertedIndex_b/StopWords.csv"
        )));
```

```
String pattern;
while ((pattern = Reader.readLine()) != null) {
    stopwordslist.add(pattern.toLowerCase());
}
```

We use a for loop in order to exclude the words from the input files that match those contained in the stopwordslist

```
for (String string : value.toString().split("\\s+")) {
    if (!stopwordslist.contains(string.toLowerCase())) {
        word.set(string.toLowerCase());
    }
}
```

The intermediate mapper output files include series of keys (word) with corresponding file(s) name where it is used. The reducer will gather all values corresponding respectively to each key.

```
context.write(word, filename);
```

In the reducer class, for each key, we add all matching values (files names in which the word is contained). All files names are separated by commas and spaces:

```
HashSet<String> set = new HashSet<String>();

for (Text value : values) {
    set.add(value.toString());
}
```



```

StringBuilder builder = new StringBuilder();

String separator = "";
for (String value : set) {
    builder.append(separator);
    separator = ", ";
    builder.append(value);
}

context.write(key, new Text(builder.toString()));

```

Below are screenshots about the execution runtime:

The first screenshot shows the Cloudera interface in a Mozilla Firefox browser. The address bar displays `http://localhost:8088/cluster/app/application_1487196663434_0005`. The left sidebar contains a 'Cluster' menu with options like 'About', 'Nodes', 'Applications', and 'Scheduler'. The main content area is titled 'Application Overview' and displays the following details:

- User:** cloudera
- Name:** InvertedIndex_b
- Application Type:** MAPREDUCE
- Application Tags:**
- State:** FINISHED
- FinalStatus:** SUCCEEDED
- Started:** Wed Feb 15 12:37:18 -1000 2017
- Elapsed:** 4mins, 0sec
- Tracking URL:** History
- Diagnostics:**

Below the overview, the 'Application Metrics' section shows:

- Total Resource Preempted:** <memory:0, vCores:0>
- Total Number of Non-AM Containers Preempted:** 0
- Total Number of AM Containers Preempted:** 0
- Resource Preempted from Current Attempt:** <memory:0, vCores:0>
- Number of Non-AM Containers Preempted from Current Attempt:** 0
- Aggregate Resource Allocation:** 1624486 MB-seconds, 1332 vcore-seconds

The 'ApplicationMaster' table shows one attempt:

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:37:18 -1000 2017	quickstart.cloudera:8042	logs

The second screenshot shows the Cloudera interface in a Mozilla Firefox browser. The address bar displays `quickstart.cloudera:19888/jobhistory/job/job_1487196663434_0005/jobhistory/job/job_1487196`. The left sidebar contains an 'Application' menu with options like 'Job', 'Overview', 'Counters', 'Configuration', 'Map tasks', and 'Reduce tasks'. The main content area is titled 'Job Overview' and displays the following details:

- Job Name:** InvertedIndex_b
- User Name:** cloudera
- Queue:** root.cloudera
- State:** SUCCEEDED
- Uberized:** false
- Submitted:** Wed Feb 15 12:37:18 HST 2017
- Started:** Wed Feb 15 12:37:40 HST 2017
- Finished:** Wed Feb 15 12:41:18 HST 2017
- Elapsed:** 3mins, 38sec
- Diagnostics:**
- Average Map Time:** 1mins, 28sec
- Average Shuffle Time:** 57sec
- Average Merge Time:** 3sec
- Average Reduce Time:** 10sec

Below the overview, the 'ApplicationMaster' table shows one attempt:

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:37:24 HST 2017	quickstart.cloudera:8042	logs

The 'Task Type' table shows the distribution of tasks:

Task Type	Total	Complete
Map	3	3
Reduce	10	10

The 'Attempt Type' table shows the distribution of attempts:

Attempt Type	Failed	Killed	Successful
Maps	0	1	3
Reduces	0	1	10

(c) (10) How many unique words exist in the document corpus (excluding stop words)? Which counter(s) reveal(s) this information? Define your own counter for the number of words appearing in a single document only. What is the value of this counter? Store the final value of this counter on a new file on HDFS.

In the code used in question b), we add the following line just before the driver class:

```
public static enum UNIQUE_WORDS_COUNTER {  
    NUMBER_OF_UNIQUE_WORDS,  
};
```

In the mapper class, stop words are excluded from the mapper output. In the reducer class, each word is added to the HashSet set. The UNIQUE_WORD_COUNTER increases by one for each unique word that appears in the document, with the following if-condition:

```
if (set.size() == 1) {  
    context.getCounter(UNIQUE_WORDS_COUNTER.NUMBER_OF_UNIQUE_WORD  
S).increment(1);  
}
```

The total number of unique words appears at the end of the terminal output. We end up with 57033 unique words:

```
[cloudera@quickstart workspace]$ hadoop jar InvertedIndex_c.jar  
stopwordcount.InvertedIndex_c input output  
[input, output]  
[input, output]  
17/02/15 12:41:53 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032  
17/02/15 12:41:57 INFO input.FileInputFormat: Total input paths to process : 3  
17/02/15 12:41:57 INFO mapreduce.JobSubmitter: number of splits:3  
17/02/15 12:41:58 INFO mapreduce.JobSubmitter: Submitting tokens for job:  
job_1487196663434_0006  
17/02/15 12:41:59 INFO impl.YarnClientImpl: Submitted application  
application_1487196663434_0006  
17/02/15 12:41:59 INFO mapreduce.Job: The url to track the job:  
http://quickstart.cloudera:8088/proxy/application_1487196663434_0006/  
17/02/15 12:41:59 INFO mapreduce.Job: Running job: job_1487196663434_0006  
17/02/15 12:42:22 INFO mapreduce.Job: Job job_1487196663434_0006 running in uber mode :  
false  
17/02/15 12:42:22 INFO mapreduce.Job: map 0% reduce 0%  
17/02/15 12:42:55 INFO mapreduce.Job: map 1% reduce 0%  
17/02/15 12:42:59 INFO mapreduce.Job: map 3% reduce 0%  
17/02/15 12:43:01 INFO mapreduce.Job: map 4% reduce 0%  
[...]  
17/02/15 12:46:07 INFO mapreduce.Job: map 100% reduce 87%  
17/02/15 12:46:10 INFO mapreduce.Job: map 100% reduce 88%  
17/02/15 12:46:14 INFO mapreduce.Job: map 100% reduce 93%  
17/02/15 12:46:18 INFO mapreduce.Job: map 100% reduce 100%  
17/02/15 12:46:21 INFO mapreduce.Job: Job job_1487196663434_0006 completed successfully  
17/02/15 12:46:22 INFO mapreduce.Job: Counters: 52  
File System Counters  
FILE: Number of bytes read=9334569  
FILE: Number of bytes written=20190356  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=25755981  
HDFS: Number of bytes written=1263324  
HDFS: Number of read operations=39  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=20  
Job Counters
```

```

Killed map tasks=1
Killed reduce tasks=2
Launched map tasks=4
Launched reduce tasks=11
Data-local map tasks=4
Total time spent by all maps in occupied slots (ms)=350410
Total time spent by all reduces in occupied slots (ms)=841723
Total time spent by all map tasks (ms)=350410
Total time spent by all reduce tasks (ms)=841723
Total vcore-seconds taken by all map tasks=350410
Total vcore-seconds taken by all reduce tasks=841723
Total megabyte-seconds taken by all map tasks=358819840
Total megabyte-seconds taken by all reduce tasks=861924352
Map-Reduce Framework
  Map input records=507536
  Map output records=507536
  Map output bytes=8319437
  Map output materialized bytes=9334689
  Input split bytes=381
  Combine input records=0
  Combine output records=0
  Reduce input groups=71037
  Reduce shuffle bytes=9334689
  Reduce input records=507536
  Reduce output records=57033
  Spilled Records=1015072
  Shuffled Maps =30
  Failed Shuffles=0
  Merged Map outputs=30
  GC time elapsed (ms)=5425
  CPU time spent (ms)=58300
  Physical memory (bytes) snapshot=1796476928
  Virtual memory (bytes) snapshot=19568939008
  Total committed heap usage (bytes)=1104359424
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=25755600
File Output Format Counters
  Bytes Written=1263324
stopwordcount.InvertedIndex_c$UNIQUE_WORDS_COUNTER
  NUMBER_OF_UNIQUE_WORDS=57033

```

Below are screenshots for the execution runtime.

Application Overview

User: cloudera
 Name: InvertedIndex_c
 Application Type: MAPREDUCE
 Application Tags:
 State: FINISHED
 FinalStatus: SUCCEEDED
 Started: Wed Feb 15 12:41:59 -1000 2017
 Elapsed: 4mins, 20sec
 Tracking URL: History
 Diagnostics:

Application Metrics

Total Resource Preempted: <memory:0, vCores:0>
 Total Number of Non-AM Containers Preempted: 0
 Total Number of AM Containers Preempted: 0
 Resource Preempted from Current Attempt: <memory:0, vCores:0>
 Number of Non-AM Containers Preempted from Current Attempt: 0
 Aggregate Resource Allocation: 1795040 MB-seconds, 1478 vcore-seconds

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:41:59 -1000 2017	quickstart.cloudera:8042	logs

Job Overview

Job Name: InvertedIndex_c
 User Name: cloudera
 Queue: root.cloudera
 State: SUCCEEDED
 Uberized: false
 Submitted: Wed Feb 15 12:41:59 HST 2017
 Started: Wed Feb 15 12:42:20 HST 2017
 Finished: Wed Feb 15 12:46:19 HST 2017
 Elapsed: 3mins, 58sec
 Diagnostics:
 Average Map Time: 1mins, 30sec
 Average Shuffle Time: 1mins, 6sec
 Average Merge Time: 4sec
 Average Reduce Time: 12sec

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:42:04 HST 2017	quickstart.cloudera:8042	logs

Task Type	Total	Complete
Map	3	3
Reduce	10	10

Attempt Type	Failed	Killed	Successful
Maps	0	1	3
Reduces	0	1	10

(d) (30) Extend the inverted index of (b), in order to keep the frequency of each word for each document. The new output should be of the form: [table on questions sheet] which means that the word frequency should follow a single '#' character, which should follow the filename, for each file that contains this word. You are required to use a Combiner.

In this question, the (key, values) formats are the following:

mapper input: (LongWritable, Text) [file names, text of the files]

mapper output – combiner input: (Text, Text) [word, file(s) name(s) where the word appears]

mapper output – combiner input: (Text, Text) [word, file(s) name(s) where the word appears]

combiner output – reducer input: (Text, Text) [word, file(s) name(s) where the word appears]

reducer output: (Text, Text) [word, file(s) name(s) where the word appears]

We reuse the code made previously and add a Combiner class. In the Reducer class, we use a HashMap that can contain a string with an integer.

```
HashMap<String, Integer> container = new HashMap<String, Integer>();
```

For each key/word, it will store files names (string) with the corresponding number of occurrences (integer). With a for-loop and an if-else condition, for each key, we span the values (file names), store them, if not already present, in the HashMap and set the corresponding integer value at 1. In case the file name was already stored, we increase the corresponding integer of the HashMap by 1.

```
for (Text value : values) {
    if (!container.containsKey(value.toString())) {
        container.put(value.toString(), 1);
    } else {
        int count = container.get(value.toString()) + 1;
        container.put(value.toString(), count);
    }
}
```

Finally, when generating the output file, for each key (word), we append the file(s) name(s) where the word appears and the number of occurrences that was stored in the Hashmap.

```
StringBuilder builder = new StringBuilder();

String separator = "";
for (String value : container.keySet()) {
    builder.append(separator);
    separator = "; ";
    builder.append(value + "#" + container.get(value));
}

context.write(key, new Text(builder.toString()));
```

Below are the screenshots with the execution time:

The image consists of two screenshots of a web browser displaying the Cloudera interface. The top screenshot shows the 'Application Overview' for a job named 'InvertedIndex_d'. The bottom screenshot shows the 'Job Overview' for the same job, including detailed task and attempt statistics.

Application Overview

User:	cloudera
Name:	InvertedIndex_d
Application Type:	MAPREDUCE
Application Tags:	
State:	FINISHED
FinalStatus:	SUCCEEDED
Started:	Wed Feb 15 12:47:19 -1000 2017
Elapsed:	5mins, 27sec
Tracking URL:	History
Diagnostics:	

Application Metrics

Total Resource Preempted:	<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:	0
Total Number of AM Containers Preempted:	0
Resource Preempted from Current Attempt:	<memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt:	0
Aggregate Resource Allocation:	2186925 MB-seconds, 1796 vcore-seconds

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:47:19 -1000 2017	quickstart.cloudera:8042	logs

Job Overview

Job Name:	InvertedIndex_d
User Name:	cloudera
Queue:	root.cloudera
State:	SUCCEEDED
Uberized:	false
Submitted:	Wed Feb 15 12:47:19 HST 2017
Started:	Wed Feb 15 12:47:53 HST 2017
Finished:	Wed Feb 15 12:52:45 HST 2017
Elapsed:	4mins, 52sec
Diagnostics:	
Average Map Time	2mins, 1sec
Average Shuffle Time	1mins, 17sec
Average Merge Time	4sec
Average Reduce Time	13sec

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Wed Feb 15 12:47:28 HST 2017	quickstart.cloudera:8042	logs

Task Statistics

Task Type	Total	Complete
Map	3	3
Reduce	10	10

Attempt Statistics

Attempt Type	Failed	Killed	Successful
Maps	0	1	3
Reduces	0	1	10

Bibliography.

White, Tom E. *Hadoop: The Definitive Guide (4th Edition)*. N.p.: O'Reilly Media, 2015.